Unified CSI plugin for All?

Experimenting Heterogeneous Storage with Single CSI Plugin for Kubernetes

Kiran Mova, Chief Architect, Co Founder, MayaData
Mohammad Asif Siddiqui, Senior System Architect, Huawei

Speakers

Kiran Mova

Chief Architect, Co-Founder
MayaData Inc
TSC, SODA Foundation

Mohammad Asif Siddiqui

Senior Systems Architect
Huawei
Member, ASF

@kiranmova

kiranmova

@asifdxtreme

Mohammad Asif Siddiqui

OpenEBS

Litmus

CLOUD NATIVE COMPUTING FOUNDATION

soda foundation

THE APACHE SOFTWARE FOUNDATION

VOLCANO

KubeEdge

Agenda

- Overview on Kubernetes and CSI
- Understanding SODA CSI Plugin
- One CSI Plugin for all - An Approach

Why Kubernetes matters?

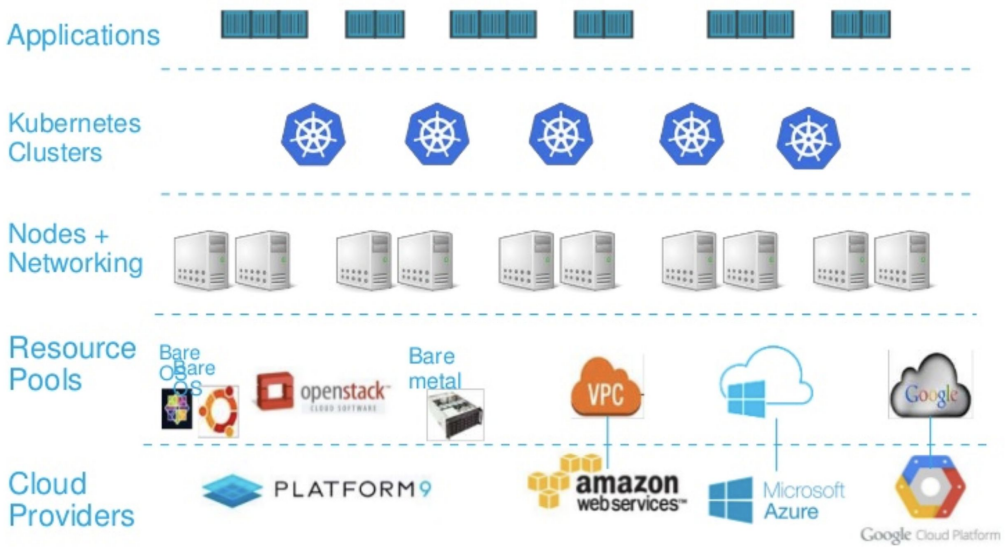

#1: Fault-tolerant by Design. Self-healing
#2: Modular App Design w/ Containers
#3: Automated Rolling Upgrades
#4: Auto-scaling. Reproducibility.
#5: Write Once, Run Anywhere. Portability.

(K8s is the new Linux)

hybrid and multi-cloud environments eases the deployment and management manages app health, replication, load balancing hardware resource allocation

CLOUD NATIVE COMPUTING FOUNDATION

https://thenewstack.io/ebooks/kubernetes/state-of-kubernetes-ecosystem-second-edition-2020/

Data on Kubernetes

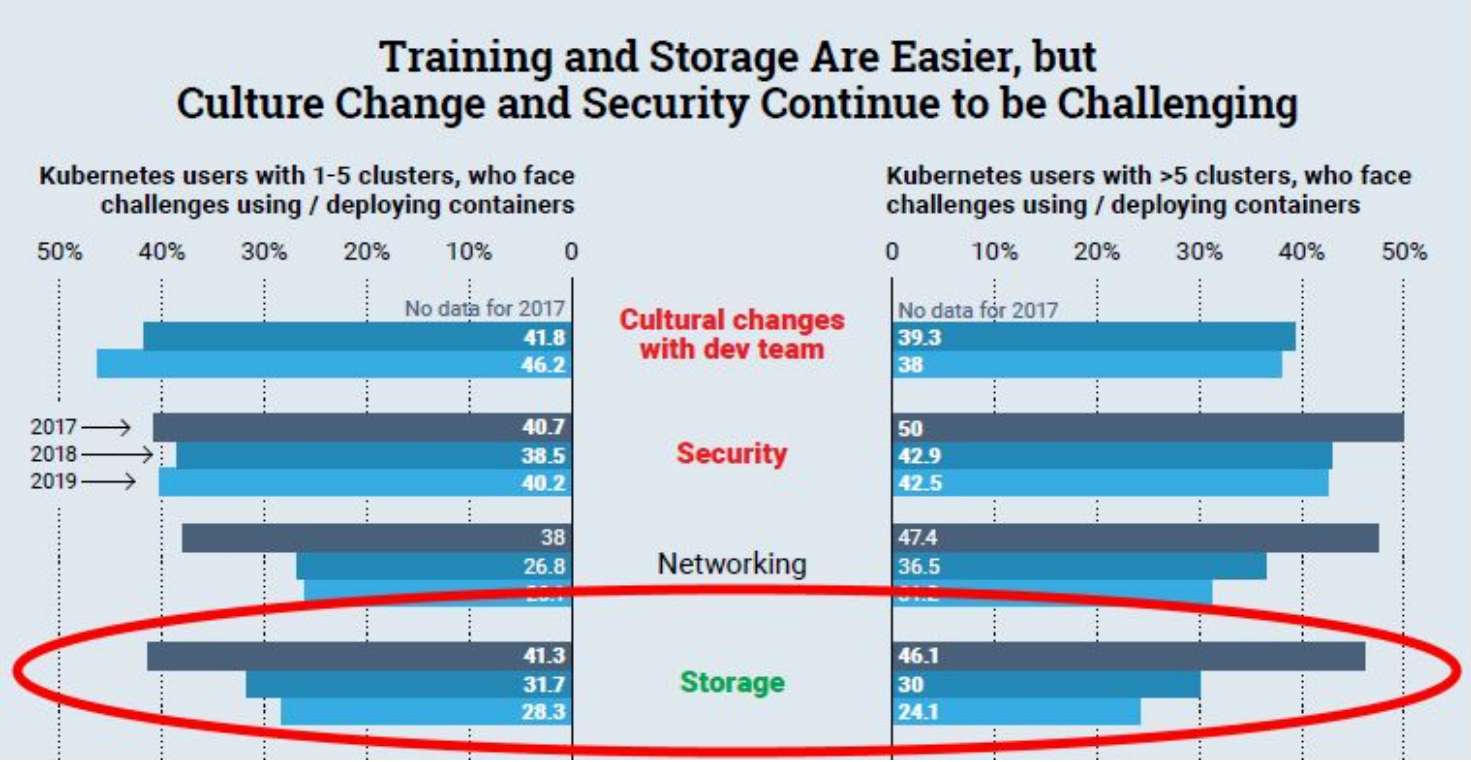

https://thenewstack.io/ebooks/kubernetes/state-of-kubernetes-ecosystem-second-edition-2020/

Data on Kubernetes using CSI

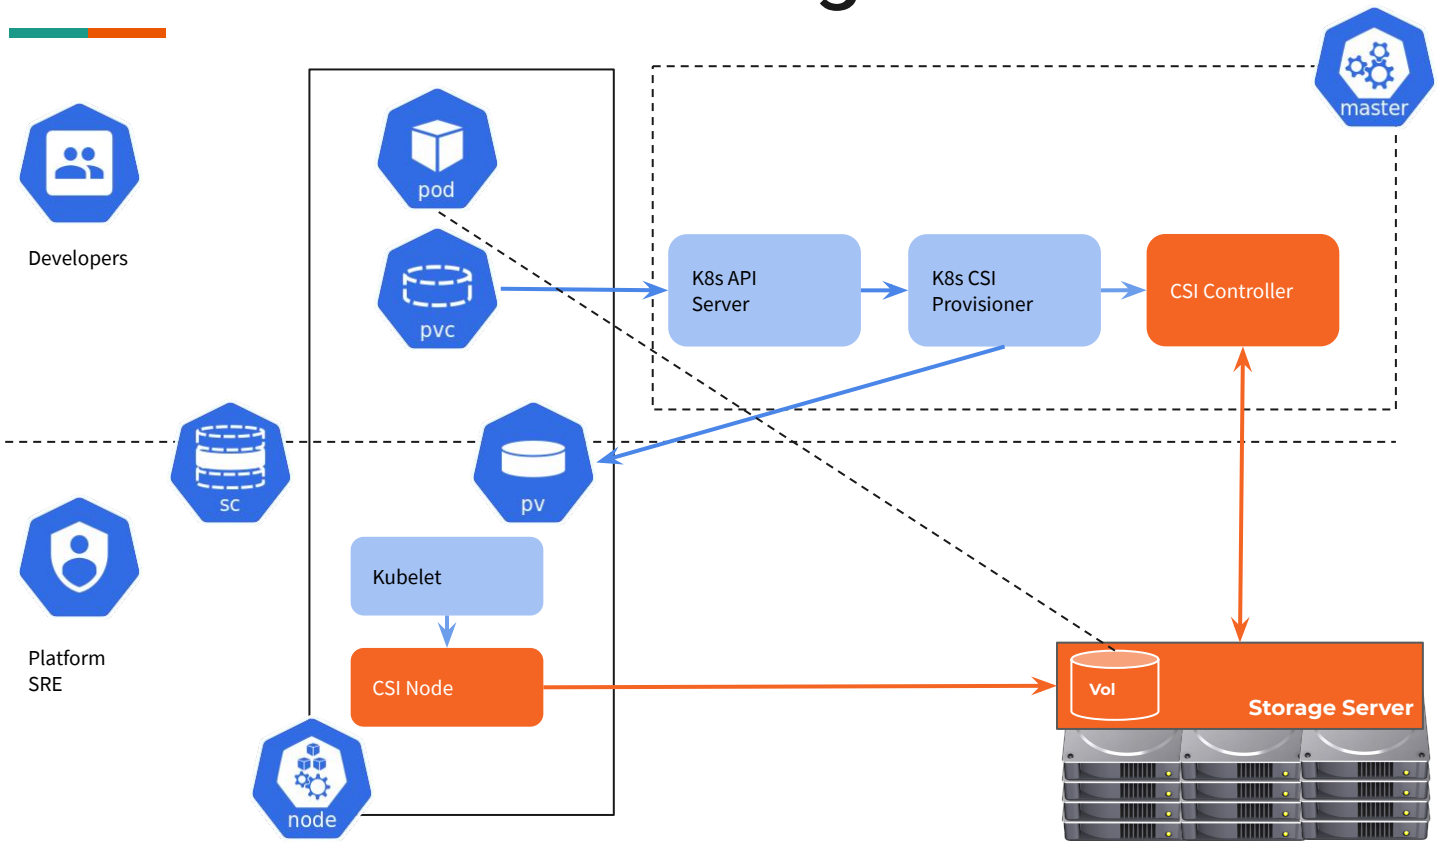

Data on Kubernetes using CSI

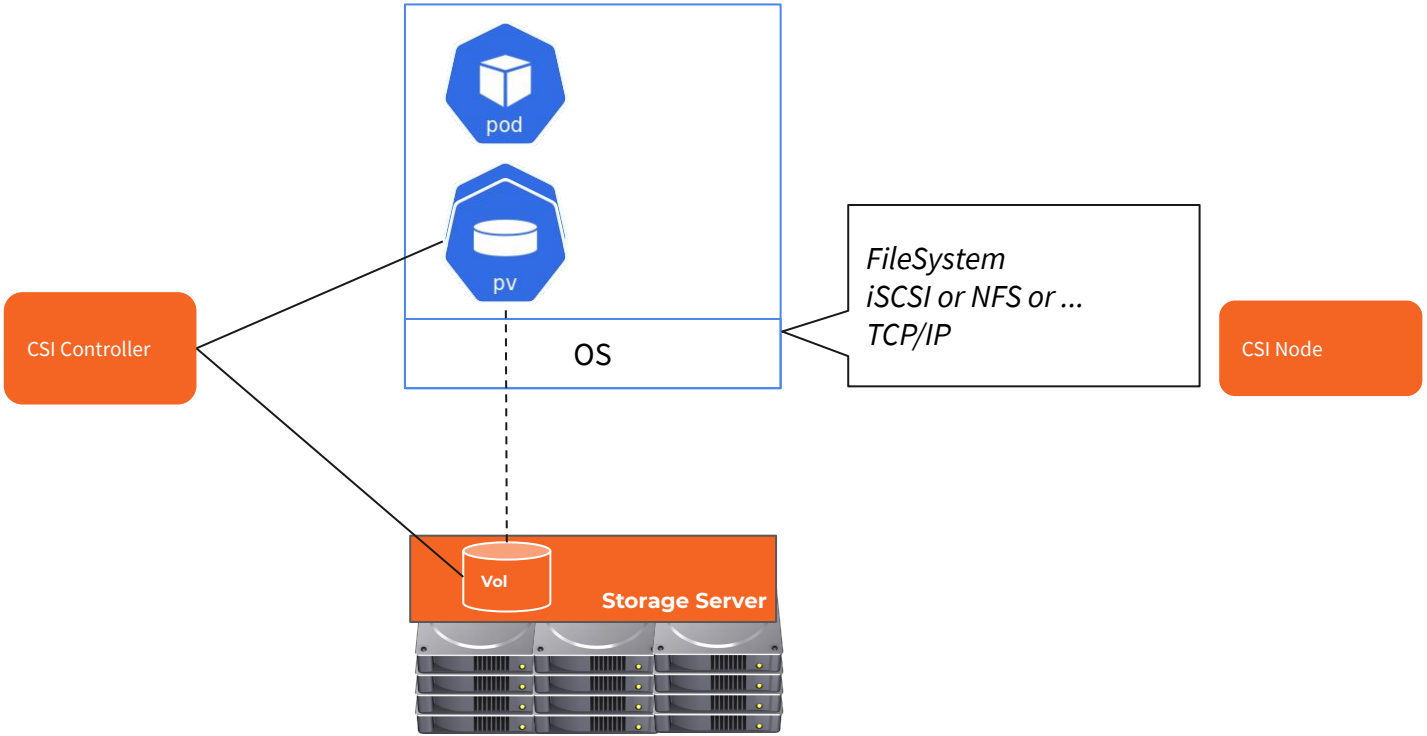

https://thenewstack.io/kubernetes-storage-dynamic-volumes-and-the-container-storage-interface/

How does CSI help end-users?

So, although ***CSI doesn't have major effects for end-users***, it correlates to a much *richer and better product offerings (90+)* from a wide range of storage vendors in the cloud native space.

- CSI solves the problem of upgrading storage functionality in Kubernetes without needing to upgrade Kubernetes itself.
- CSI thus requires that you have independent units running in a cluster which are privileged and able to speak the CSI language for attaching storage to a mount point.
- Container native storage does not require CSI in any way — you can run storage in a container and create volumes just as easily without CSI as with it. The only difference is that CSI allows vendors to have complete control over their release cycle.

https://thenewstack.io/kubernetes-storage-dynamic-volumes-and-the-container-storage-interface/

End-users: Evolving Data Infrastructure

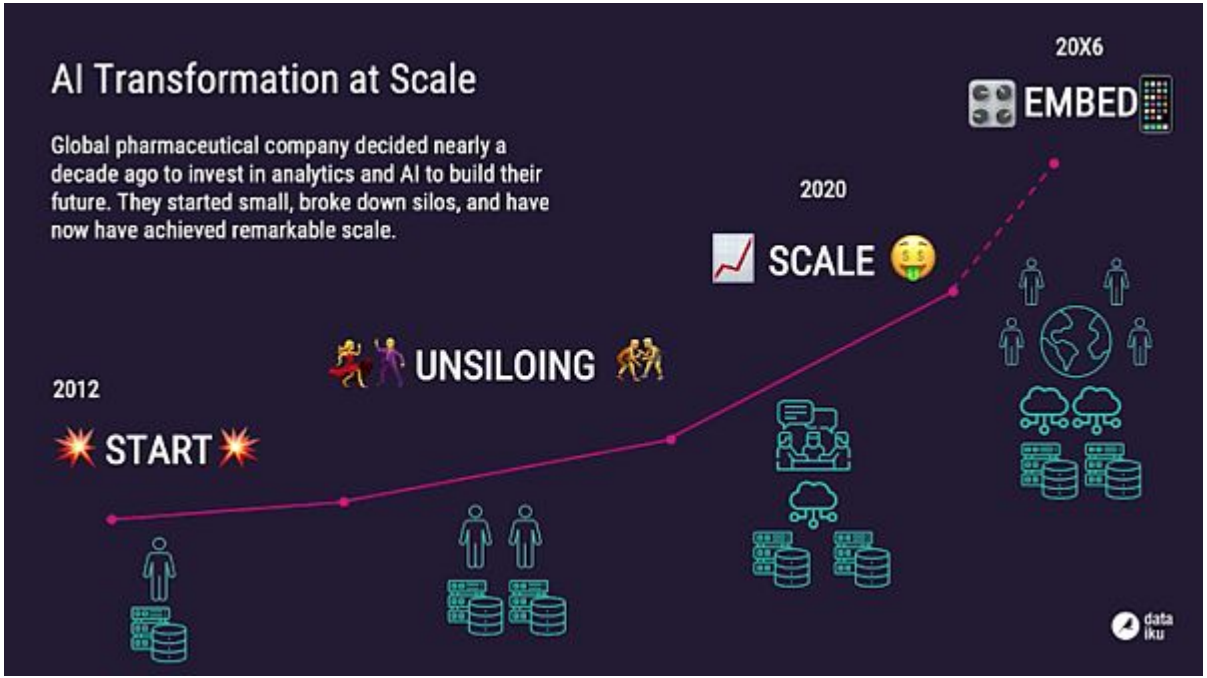

https://mattturck.com/data2020/

2020 - Big Data and AI Landscape

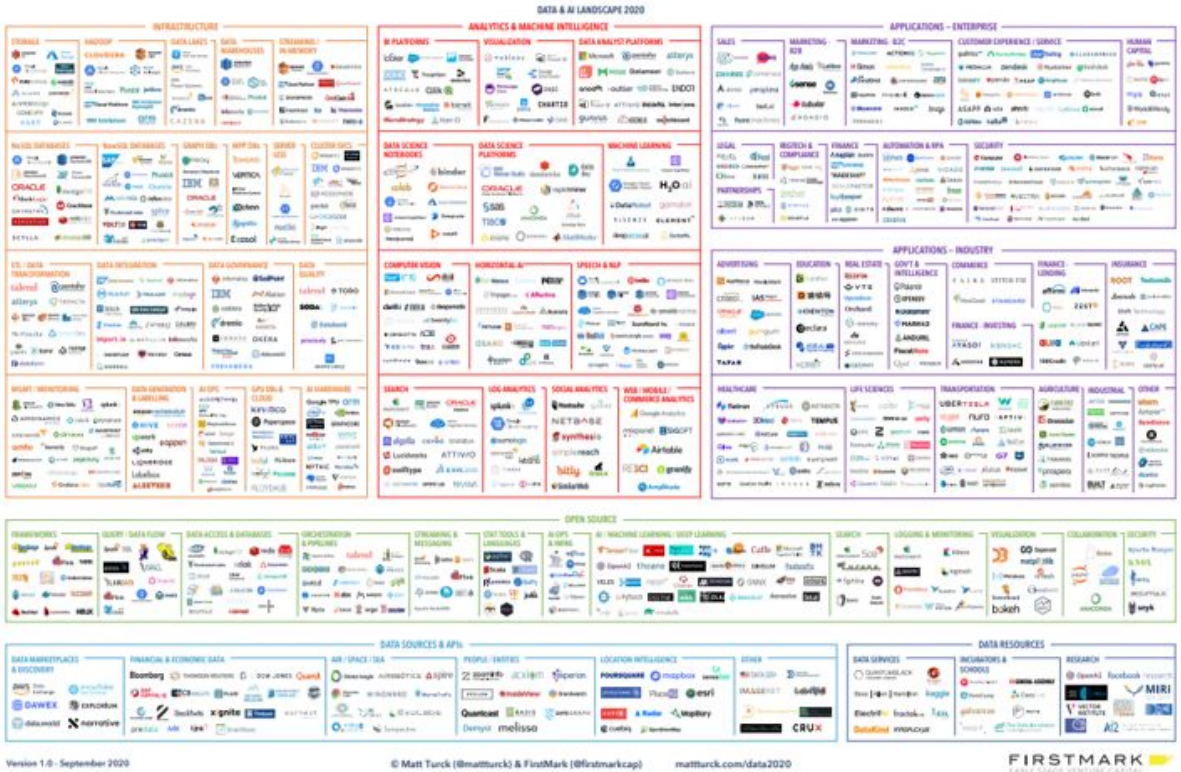

https://mattturck.com/data2020/

*"to succeed, every modern company will need to be **not just a software company, but also a data company**."*

"overall volume of data flowing through the enterprise continues to ***grow an explosive pace.*** The number of sources of data keeps increasing as well, with ever more SaaS tools. There is ***not one but many data pipelines*** operating in parallel in the enterprise."

"raises the bar on data infrastructure (and the teams building/maintaining it)"

Typical Data/Storage Stack

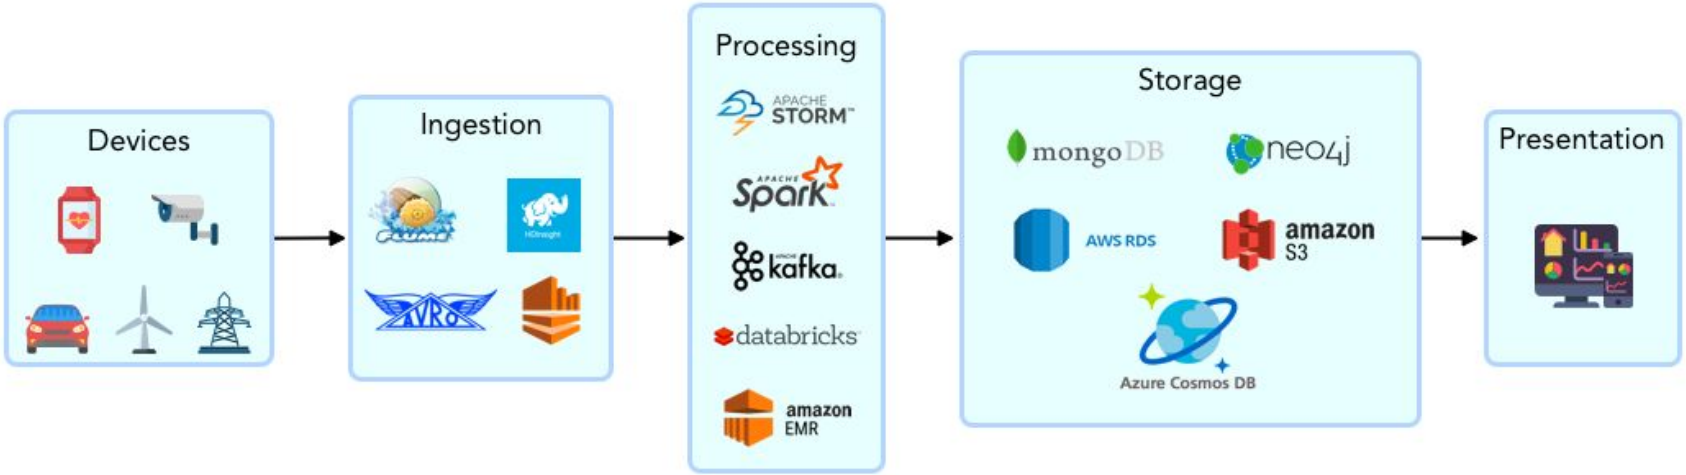

https://blog.qburst.com/2019/02/data-engineering-iot/

Workload (Stateful App) Needs

Workload	Needs Durability	Needs to be fast	Avg size per container	Dedicated disk?
AI Workloads with model or training caches	No	No	1 to 100GB	No
Postgres workloads with millions of rows	Yes	Yes	1G to 5TB	Yes
HDFS datanode backend folders	No	Yes	10TB	Yes
Cold storage	Yes	No	Infinite (assume that over time the amount of storage increases indefinitely)	No

https://thenewstack.io/kubernetes-storage-dynamic-volumes-and-the-container-storage-interface/

Persistent Volume types

External Volumes *(SAN/NAS)*

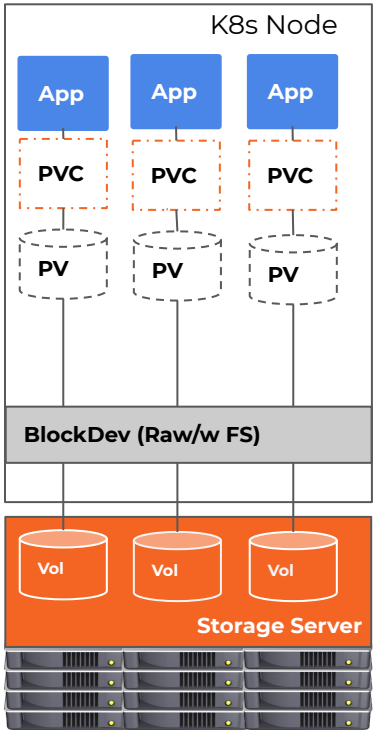

(DAS) Local Volumes

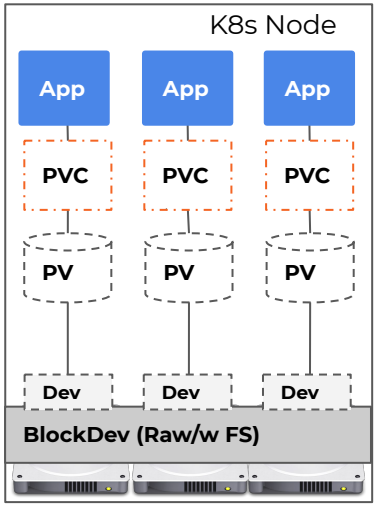

CAS Volumes

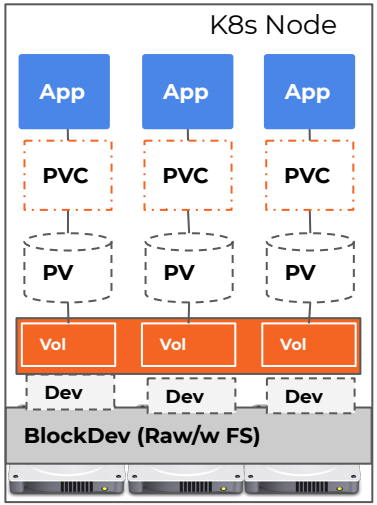

NAS - Network Attached Storage *(Example: GPD, EBS, Storage Appliances)*
DAS - Direct Attached Storage *(Example: hostDir, Local PV)*
CAS - Container Attached Storage *(Example: OpenEBS)*

Represent stateful Pods like Databases, etc.

Indicates functionality like replication, snapshots, encryption, compression, etc.

https://www.cncf.io/blog/2020/09/22/container-attached-storage-is-cloud-native-storage-cas/

Data on Kubernetes

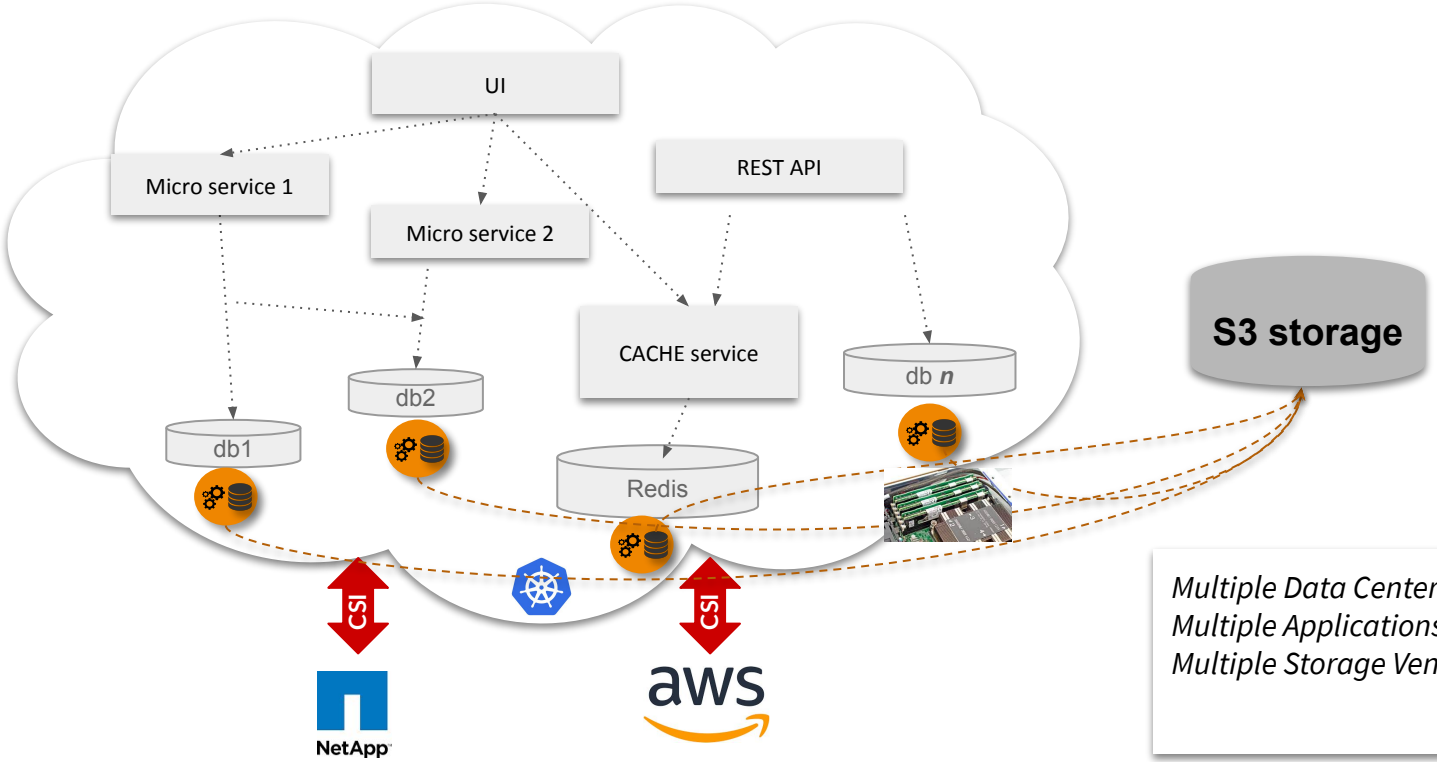

Multiple Data Centers
Multiple Applications
Multiple Storage Vendors

Data on Kubernetes using *multiple* CSI

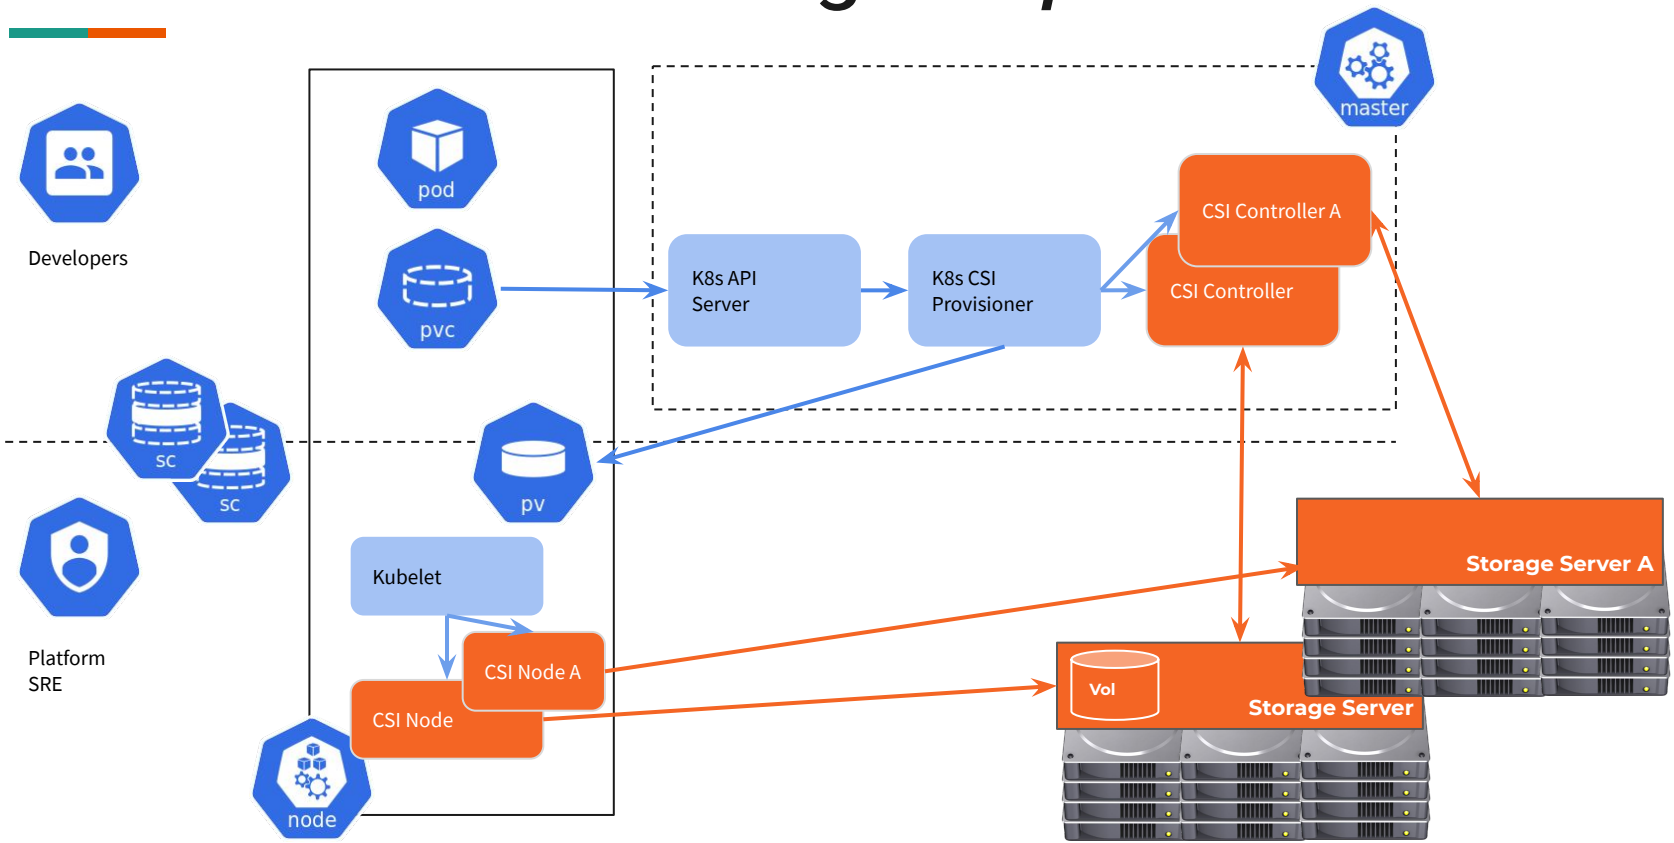

CSI Challenges with Heterogeneous Storage

Developers

All I care is about PVC. Now I have to deal with selecting which Storage Class to pick and what features it provides.

Additional Kubernetes Operators to simplify this complexity?

Platform SRE

Overhead of maintaining life cycle (install and upgrades) of multiple CSI drivers. The complexity of managing the compatibility between Kubernetes versions and CSI Drivers.

More pods! Overhead on system resources.

In-tree volumes were better!

Evolving Specification. (*Data Protection!)*

Doesn't deal with Data Migration across multiple providers or moving across multiple Kubernetes Clusters.

Doesn't manage the underlying storage, which comes with its own challenges even for a single provider. *How to deal with failed drives? How to ensure availability zones?*

And VMs and Bare Metal!

soda foundation

Who provides storage to K8s itself?

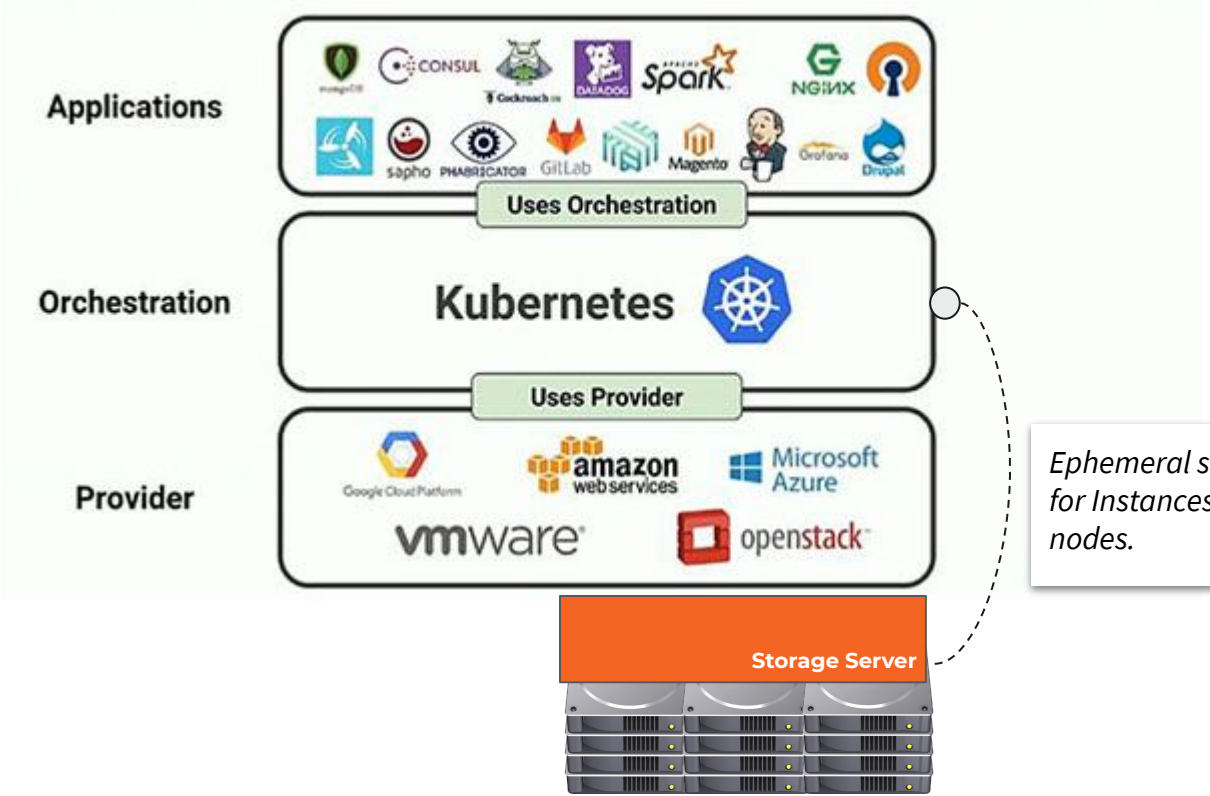

Ephemeral storage (aka Root Disks) for Instances running Kubernetes nodes.

soda foundation

Data lives on, K8s clusters don't

Cluster Lifecycle

Augment K8s capabilities that can be used by Data(Application) Operators.

DATA SCHEDULERS

DATA SECURITY

Data Operators

Data Lifecycle

Manage the data beyond the life of a single K8s cluster. And beyond the application purview.

DATA PROTECTION

DATA MIGRATION

DATA GOVERNANCE

DATA COMPLIANCE

UNIFIED STORAGE PLATFORM

soda foundation

SODA: Unified CSI Driver?

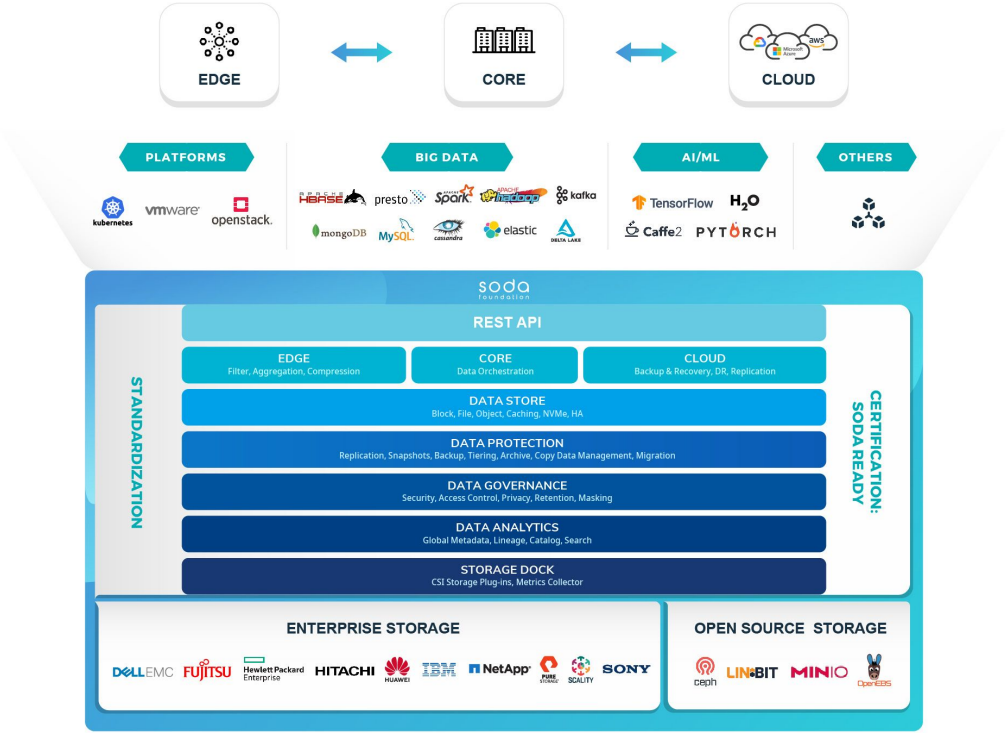

So, although *CSI does help storage vendors with Kubernetes integration*, ***CSI doesn't solve the growing use cases for Data management.***

soda foundation

Soda ecosystem and it's CSI Plugin

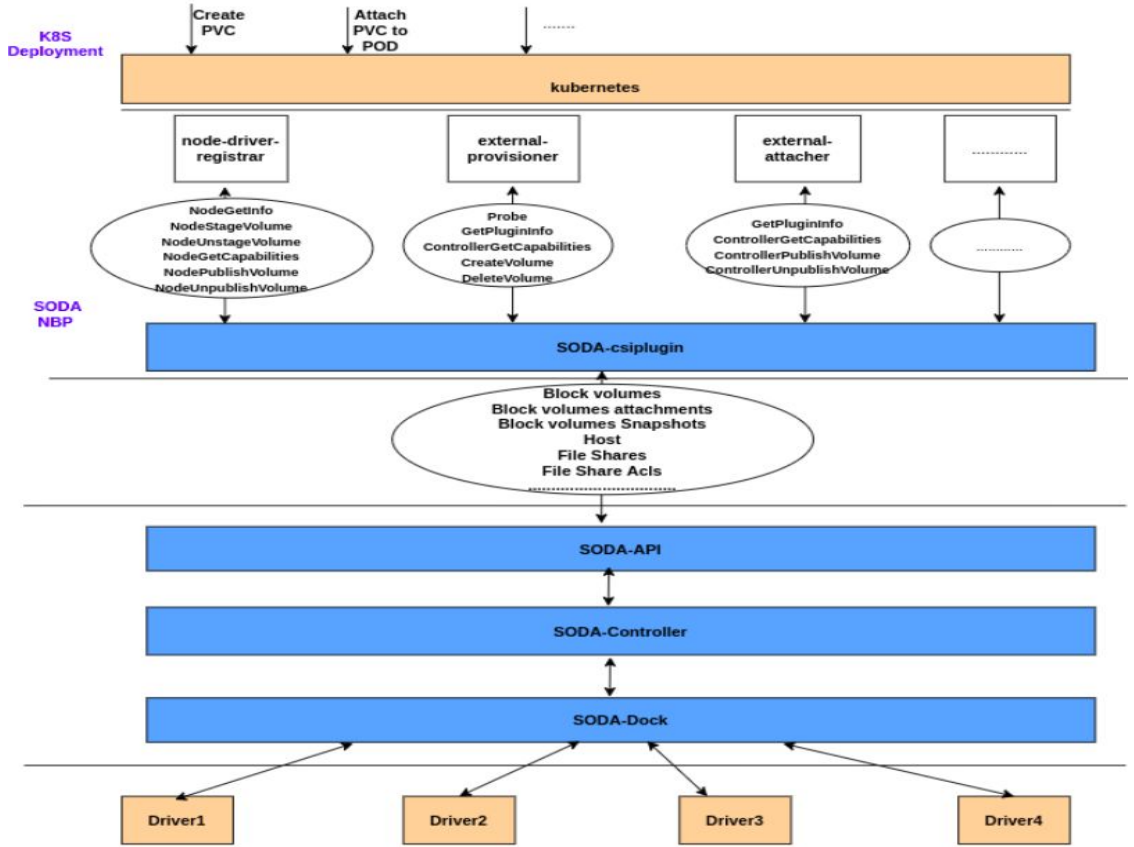

CSI PlugNPlay (One CSI plugin for all heterogenous drivers)

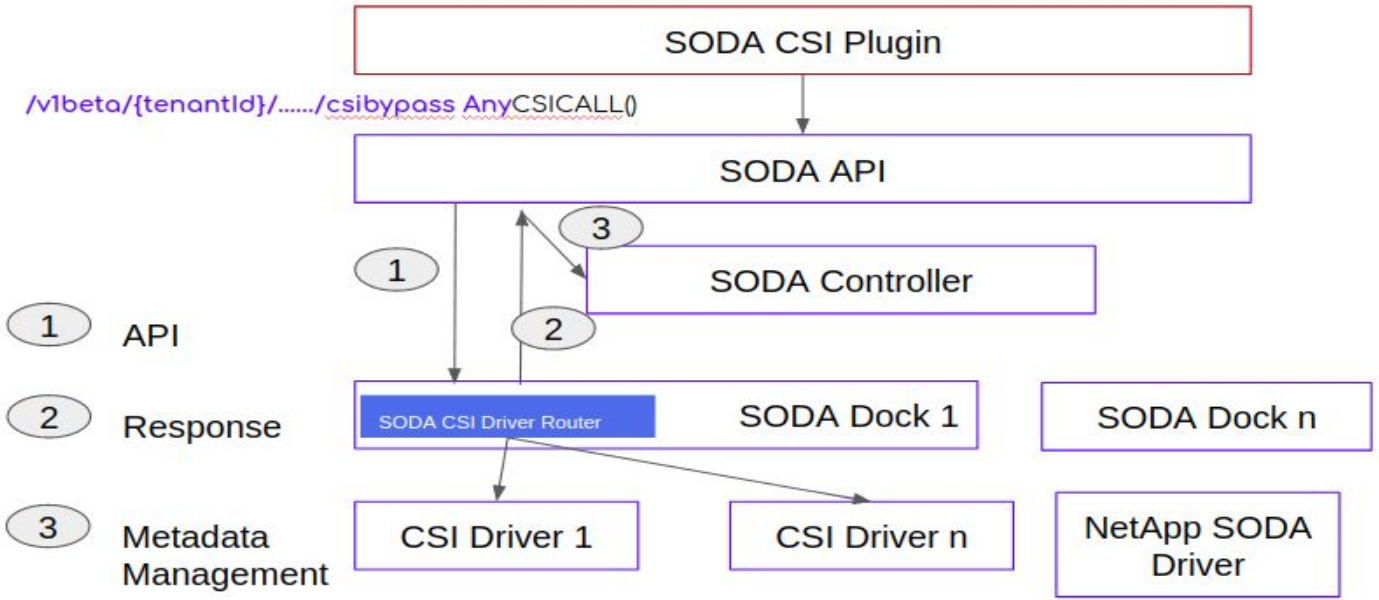

Volume Provisioning for PlugNPlay

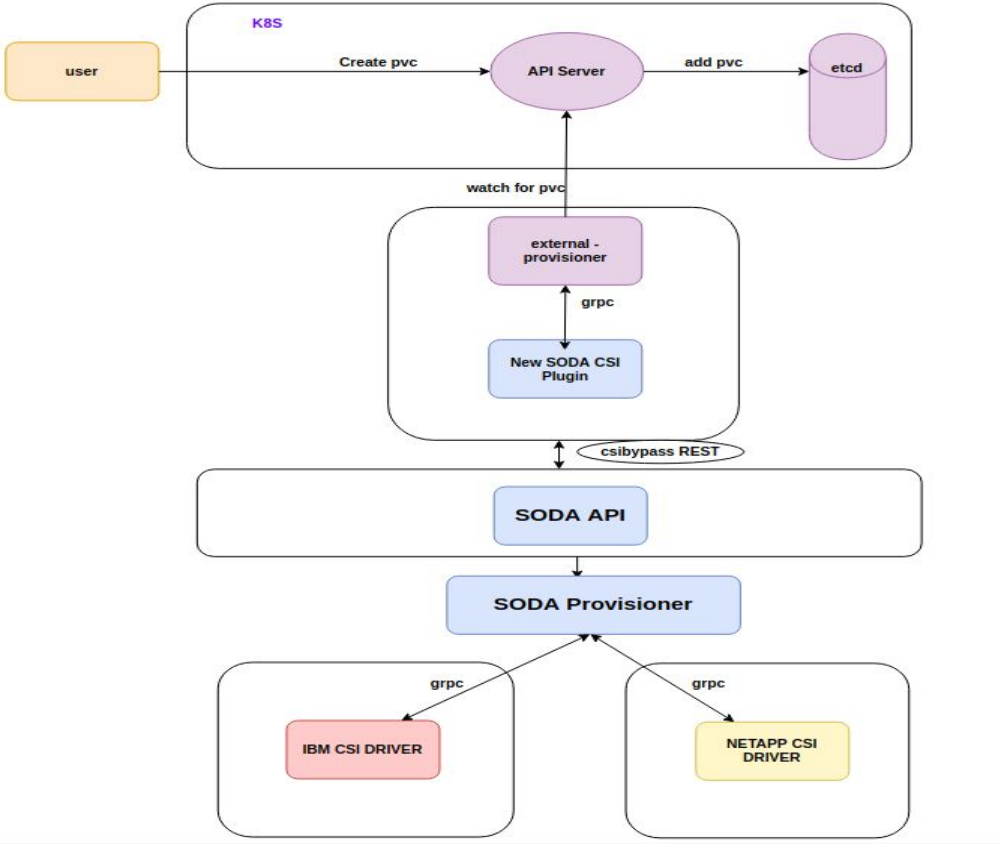

Soda CSI PlugNPlay (Alternative Approach)

Goals

Make a single CSI Plug-N-Play mechanism which helps to:

- Provision and manage the Storage of heterogeneous Storage providers.
- Third party CSI drivers should be used directly from Vendors so that SODA/Users need not worry about it's maintenance.
- Use Soda Profile ID to determine which storage vendor drivers need to be used to provision the storage.
- Experience all the features set of SODA.

Deployment View

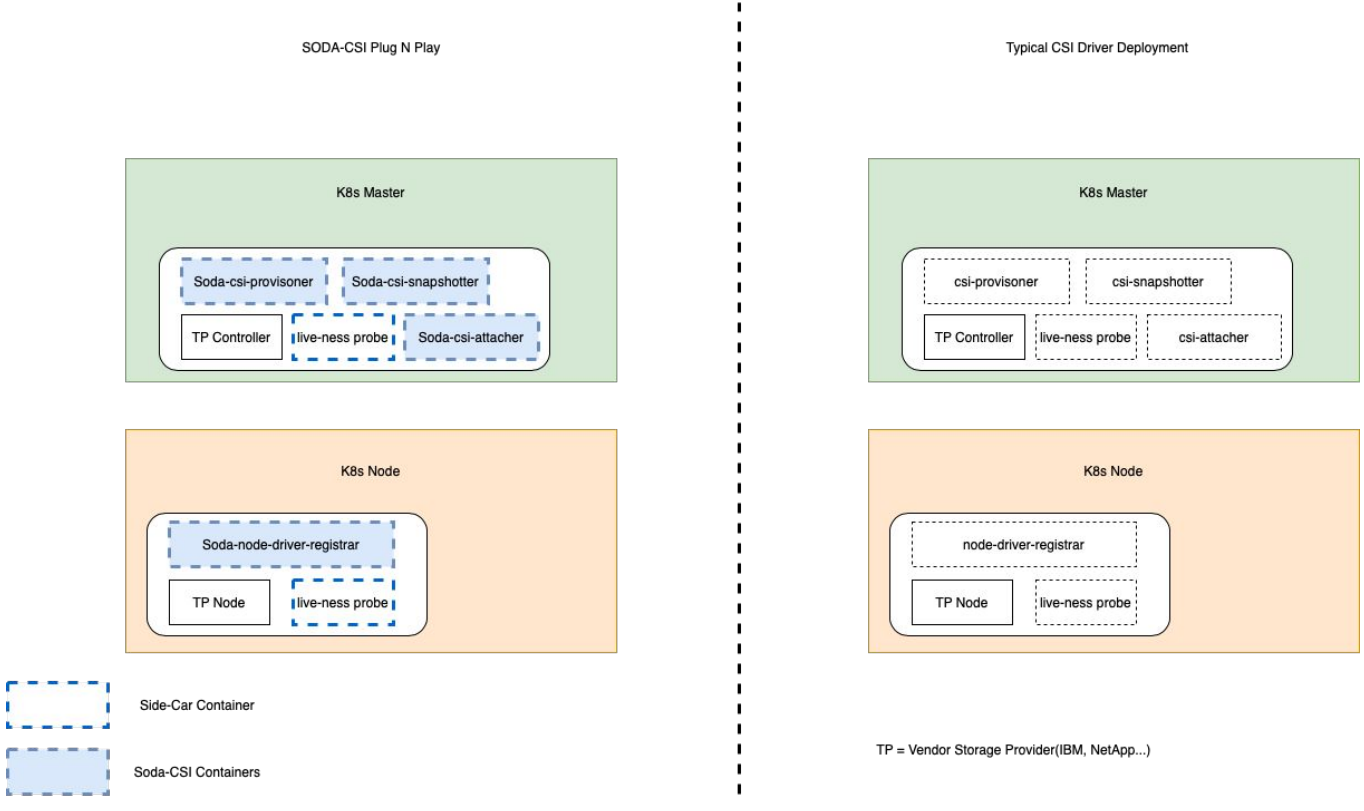

Call flows

A typical call flow for SODA CSI Plug-N-Play mechanism will be:

- Once the PVC Object is created by user then soda-csi-provisioner will get a request to provision the Storage.
- soda-csi-provisioner will interact with SODA API's to get the profile details and other parameters to determine if the required driver is same as the current driver in a pod.
- soda-csi-provisioner updates the SODA API server with the volume provisioning request and gather the required intelligence from SODA about the current provisioning.
- soda-csi-provisioner will forward the request to CSI driver in the same pod to do the actual provisioning.

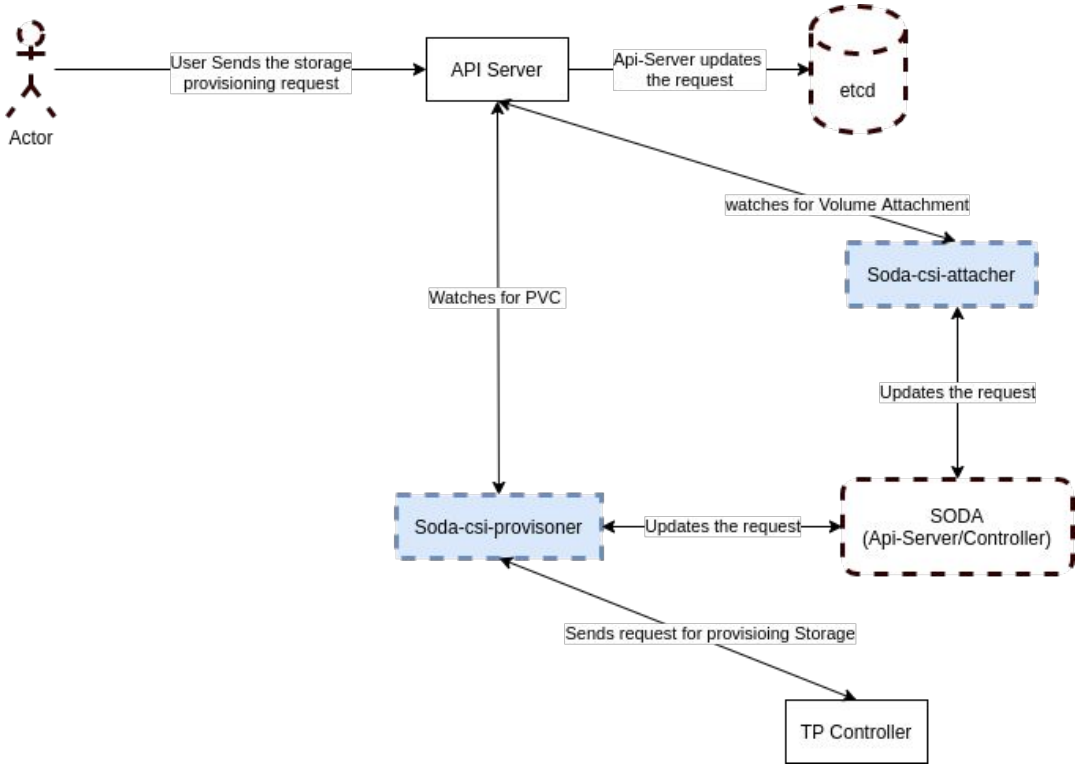

Advantages of the SODA CSI PlugNPlay

- A simple design which let's user to patch the existing Third party CSI drivers deployment without any changes to actual CSI driver containers provided by TP vendors.
- All the sidecar containers which will be used by TP CSI drivers will be provided by SODA.
- Robust and easy to maintain design which can be used by heterogeneous csi storage providers with minimal changes to side cars.
- Easy to upgrade along with TP CSI driver.
- Easy to maintain the side car code as the code-base will be used from kubernetes-csi org and the SODA features will be added to it as a plugin, so it's easy to upgrade side-car as k8s csi spec evolves.

Demo Video

Q and A

Additional Information

- SODA slack. https://sodafoundation.io/slack/
- https://github.com/sodafoundation/design-specs
- https://www.meetup.com/OpenDataAutonomy/
- https://www.meetup.com/Bangalore-CNCF-Meetup/
- https://www.meetup.com/Cloud-Native-Data-Management/
- https://www.meetup.com/Data-on-Kubernetes-community/
- Demo Video : https://youtu.be/ytXY_dKQCYg


@kiranmova


@asifdxtreme


kiranmova


Mohammad Asif Siddiqui

Kubernetes and CSI Interaction

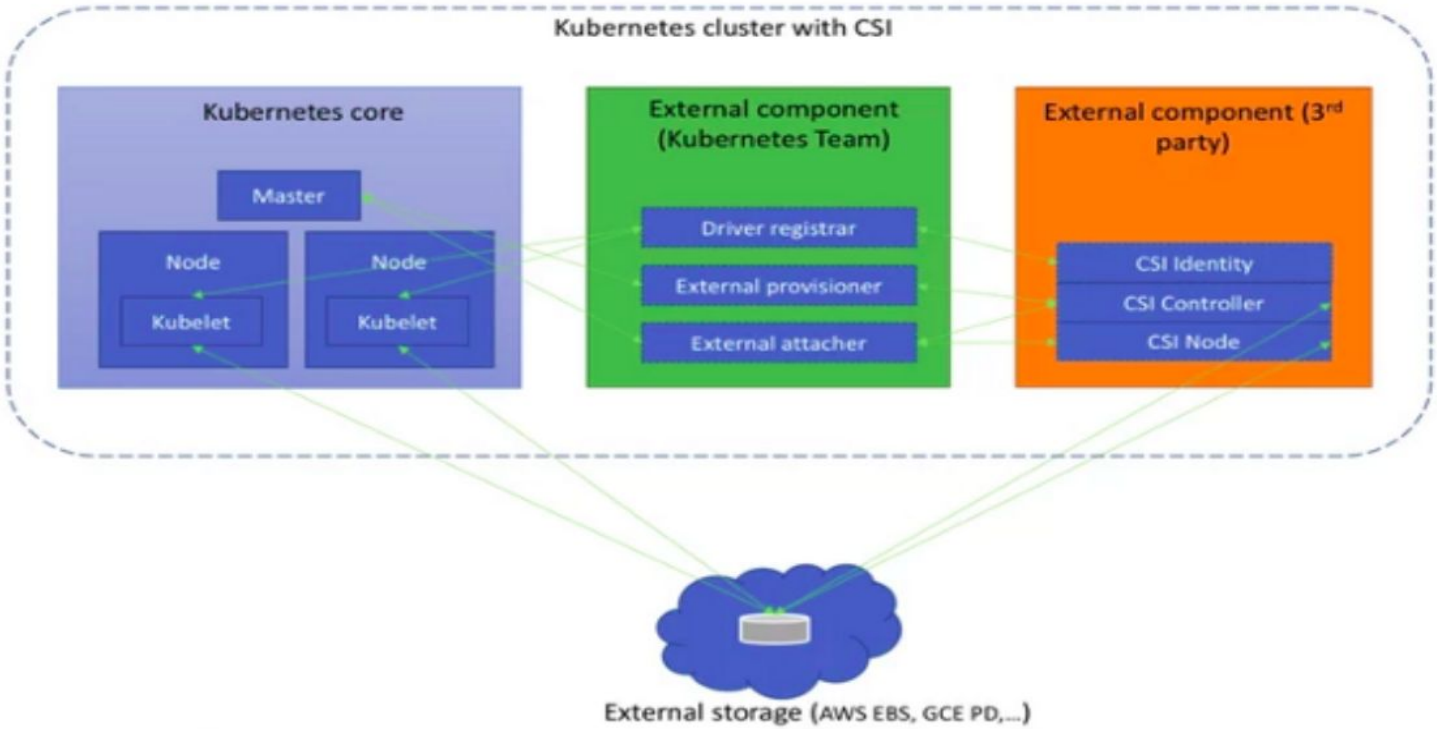

Normal Provisioning for CSI Drivers