# SOCION 2020 DATA CONNECTED

#sodacon2020

Spark on Kubernetes - Challenges in Storage & Compute Layer

Sunil Govindan / 10th Dec 2020



## SPEAKER BIO



**SUNIL GOVINDAN** 

Engineering Manager @ Cloudera

Apache Hadoop, Submarine, YuniKorn PMC Member & Committer

@sunilgovind

www.linkedin.com/in/sunilgovindan





## Agenda



### Why Spark on K8s

Compute: Gaps of K8s Scheduling for Spark

Spark Scheduling with YuniKorn & Features Overview

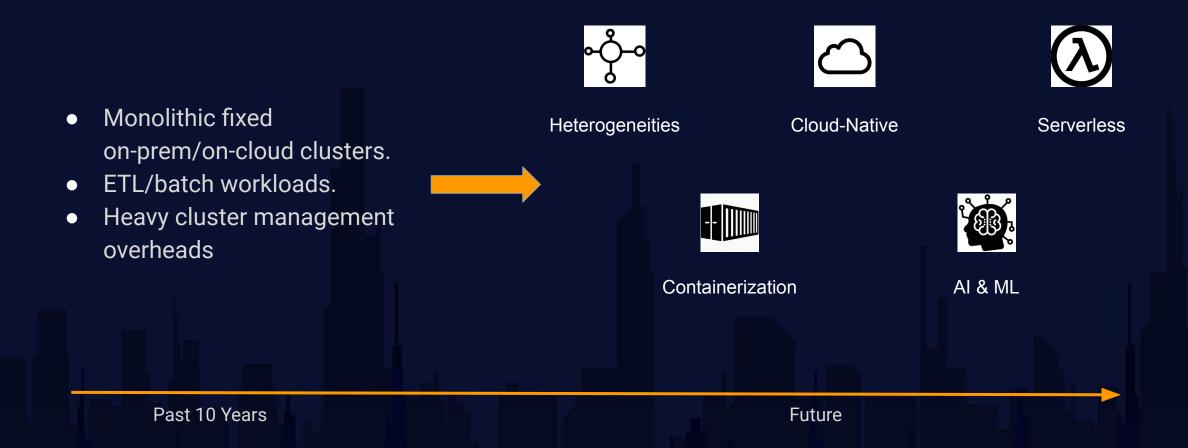
Storage: Challenges in K8s

Common Storage Layer - Introduction

# WHY SPARK ON K8s

## TODAY'S RESOURCE MANAGEMENT LANDSCAPE





## TODAY'S RESOURCE MANAGEMENT LANDSCAPE



• Kubernetes helps a lot of



Cloud-Native



Containerization



AI & ML

Industry is started to move batch to Kubernetes.

## ABOUT K8s

# The de facto container orchestration system

- Good Isolation (runtime, network and storage)
- Highly pluggable/extensible (CRD, operators, plugins)
- Declarative Language
- Loose coupled components





### WHY CHOOSE K8s TO RUN SPARK



- Containerized spark compute to provide shared resources across different ML and ETL jobs
- Support for multiple Spark versions, Python versions, and version controlled containers on the shared K8s clusters for both faster iteration and stable production
- A single, unified infrastructure for both majority of our data compute and micro services with advanced, unified observability and resource isolation support
- Fine grained access controls on shared clusters

Why Spark on K8s

**Compute: Gaps of K8s Scheduling for Spark** 

Spark Scheduling with YuniKorn & Features Overview

Storage: Challenges in K8s

Common Storage Layer - Introduction



## GAPS IN K8s SCHEDULING

### Compute in Big Data world is complex

- Long running & stateful services, HBase, Hive-LLAP, Flink, Kafka...
- Batch workloads, MR, Spark, Tez, Tensorflow ...

### Gaps in scheduling:

- Lack of 1st class application concept
- Lack of fine-grained capacity management
- App/task gets allocation in disorder
- NO resource fairness between tenants
- Throughput, scheduling latency ...

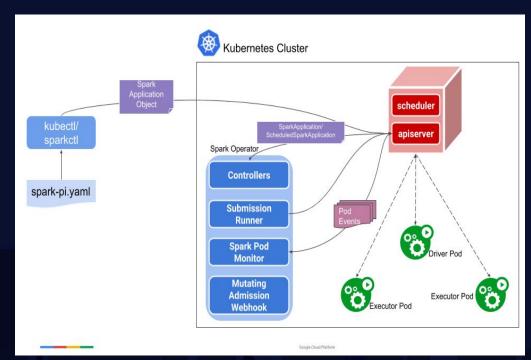




## SPARK on K8s: Deployment Flavours

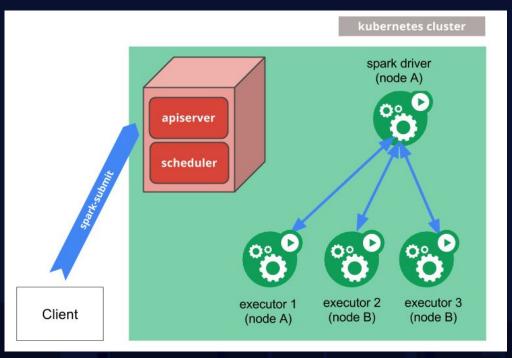


### Native Spark on K8s



Identify Spark jobs by the pod labels

### Spark K8s Operator



Identify Spark jobs by CRDs (e.g SparkApplication)







### **Job Scheduling Requirements**

- Job ordering/queueing
- Job level priority
- Resource fairness (between jobs / queues)
- Gang scheduling

### **Resource Sharing and Utilization Challenges**

- Spark driver pods occupied all resources in a namespace
- Resource competition, deadlock between large jobs
- Misbehave jobs could abuse resources

K8s default scheduler was NOT created to tackle these challenges





- High latency (~100 seconds) using the default scheduler is observed on a single K8s cluster for large volumes of batch workloads
- Large batch fair sharing in the same resource pool is unpredictable with the default scheduler
- The need for an elastic and hierarchical priority management for jobs in K8s
- Richer and online user visibility into the scheduling behavior
- Simplified layers of controllers with custom K8s scheduler



Why Spark on K8s

Compute: Gaps of K8s Scheduling for Spark

**Spark Scheduling with YuniKorn & Features Overview** 

Storage: Challenges in K8s

Common Storage Layer - Introduction



## APACHE YUNIKORN (INCUBATING)

# • •/•

### What is it:

- A standalone resource scheduler for K8s
- Focus on building scheduling capabilities to empower Big Data on K8s

## Simple to use:

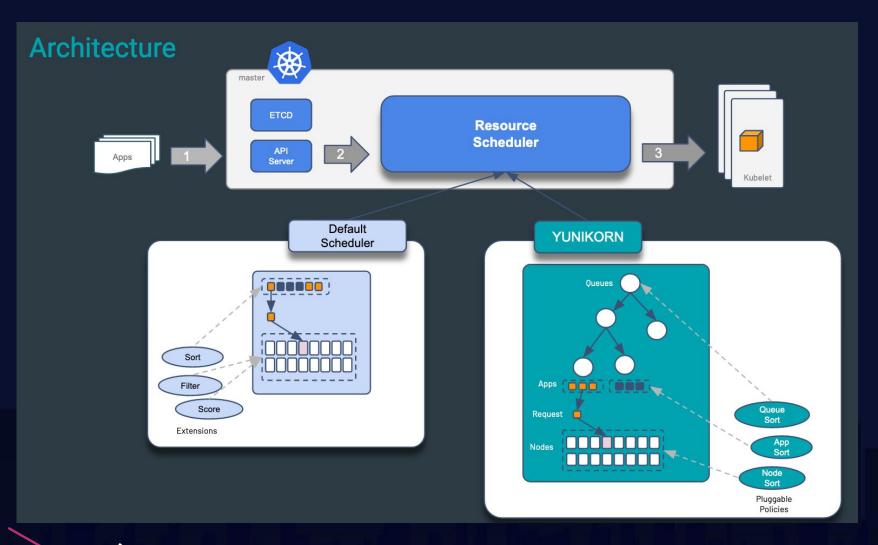
- A stateless service on K8s
- A secondary K8s scheduler or replacement of the default scheduler





## SCHEDULING WITH YUNIKORN





YuniKorn QUEUE, APP concepts are critical to provide advanced job scheduling and fine-grained resource management capabilities





Feature	Default Scheduler	YUNIKORN	Note
Application concept	×	V	Application is the 1st class citizen in YuniKorn. YuniKorn schedules apps with respect to, e,g their submission order, priority, resource usage, etc.
Job ordering	×	1	YuniKorn supports FIFO/FAIR/Priority (WIP) job ordering policies
Fine-grained resource capacity management	X		Manage cluster resources with hierarchy queues. Queues provide the guaranteed resources (min) and the resource quota limit (max).
Resource fairness	X	<b>√</b>	Resource fairness across application and queues to get ideal allocation for all applications running
Natively support Big Data workloads	×		Default scheduler focuses for long-running services. YuniKorn is designed for Big Data app workloads, and it natively supports to run Spark/Flink/Tensorflow, etc efficiently in K8s.
Scale & Performance	X	Samuel Samuel	YuniKorn is optimized for performance, it is suitable for high throughput and large scale environments.



## **RUNNING SPARK WITH YUNIKORN**





### Submit a Spark job

- 1) Run spark-submit
- 2) Create SparkApplication CRD

Job is Starting

Spark driver is running

Spark executors are created

Spark job is running



Spark Driver pod

Api-server creates the driver pod

Spark Driver pod

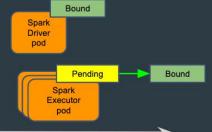
Api-server binds the pod to the assigned node



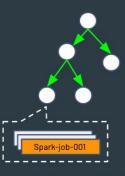
Driver pod is started, it requests for Spark executor pods from api-server



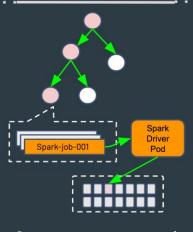
Driver pod requests for executors, api-server creates executor pods



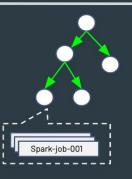




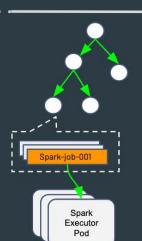
Spark job is registered to YuniKorn in a leaf queue



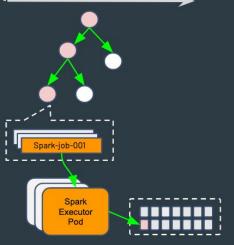
Sort queues -> sort apps -> select request -> select node



Ask api-server to bind the pod to the node



New executors are added as pending requests



Schedule, and bind executors

SCHEDULER THROUGHPUT

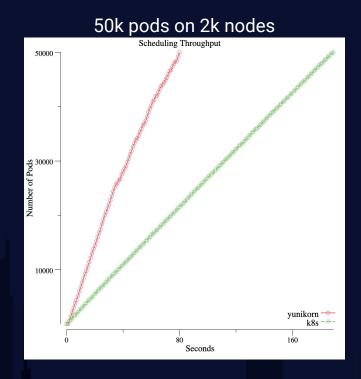
Schedule 50,000 pods on 2,000/4,000 nodes.

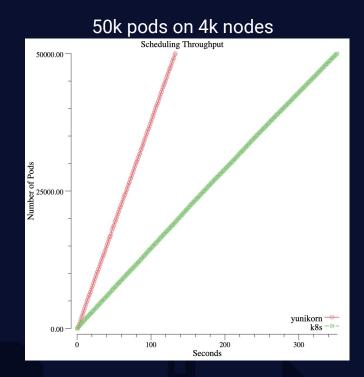
**Compare Scheduling** throughput (Pods per second allocated by scheduler)

Red line (YuniKorn) **Green line** (Default Scheduler)

**↑ 134%** 617 vs 263

**↑ 164%** 373 vs 141





### Detail report:

https://github.com/apache/incubator-yunikorn-core/blob/master/docs/evaluate-perf-function-with-Kubemark.md



### **K8s COMPATIBLE**

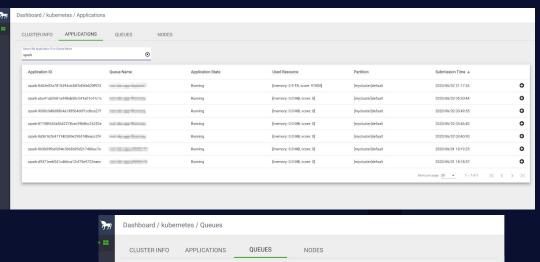
- Support K8s Predicates
  - Node selector
  - Pod affinity/anti-affinity
  - Taints and toleration
- Support PersistentVolumeClaim and PersistentVolume
  - Volume bindings
  - Dynamical provisioning
- Publishes key scheduling events to K8s event system
- Work with cluster autoscaler
- Support management commands
  - cordon nodes

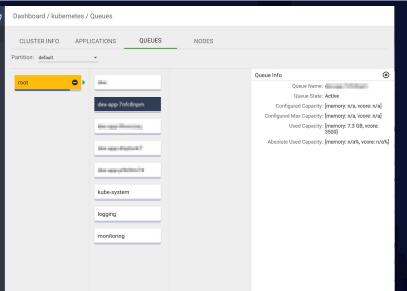


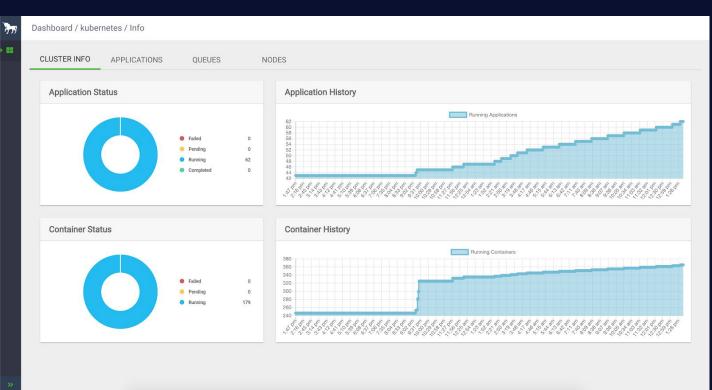












### **ROADMAP**

Open source Apache release: planned every 3 months

Next release main features (v0.10, October/November):

- Kubernetes upgrade to 1.16
- Logging and tracing enhancement using OpenTracing
- Application tracking using Kubernetes CRD
- Application and task priority support
- Web UI and REST API enhancements



### COMMUNITY





- Deployed in non-production K8s clusters
- Launched 100s of large jobs per day on some of the YuniKorn queues
- Reduced our large job scheduler
  latency by factor of ~ 3x at peak time
- K8s cluster overall resource utilization efficiency (cost per compute) improved over the default kube-scheduler for mixed workloads
- FIFO and FAIR requests are more frequently met than before

### **CLOUDERA**

- Shipping with Cloudera Public
  Cloud offerings
- Provide resource quota management and advanced job scheduling capabilities for Spark
- Responsible for both micro-service, and batch jobs scheduling
- Running on Cloud with auto-scaling enabled

## Alibaba.com<sup>®</sup>

- Deployed on pre-production on-prem cluster with ~100 nodes
- Plan to deploy on 1000+ nodes production K8s cluster this year
- Leverage YuniKorn features such as hiercharchy queues, resource fairness to run large scale Flink jobs on K8s
- Gained x4 scheduling performance improvements



Why Spark on K8s

Compute: Gaps of K8s Scheduling for Spark

Spark Scheduling with YuniKorn & Features Overview

**Storage: Challenges in K8s** 

Common Storage Layer - Introduction



## Challenges in Storage Side

Spark requires data access in various dimensions like other Big Data applications. Other than normal data input and output, Spark also requires storage requirements for intermediate data.

Some of the known challenges in this space are,

- Lack of a reliable external shuffle handler service.
- Lack of universal storage access for various storage types like ceph, portworx etc.
- Multi-cloud data access for data input/output. Could S3 be a standard way to read/write data across?



## Challenges in Storage Side: External Shuffle

\_\_\_

Shuffle data is a cause of concern for Spark applications in K8s. Unlike in YARN, there are no external shuffle service to optimize this.

- Persistent Volume is one option to overcomes this, however this is found costlier and slower.
- Remote Shuffle is another option, however this requires lot of efforts to design and manage the shuffle service outside K8s.

## Challenges in Storage Side: Multiple Storage Providers - On-prem



Many of the big-data use case requires strong Persistent Volume Support.

Few of the choices are

- Rook Ceph
- Portworx
- OCS

It is difficult to manage and maintain various types of storage types!



## Challenges in Storage Side: Multiple Storage Providers - On-prem



### For end user,

- Could users need to choose the PV Storage Class for each deployment?
- Could some of the storage class be better than other for different use cases w.r.t performance?
- Could certain storage type handle data for faster access, and certain other storage type be better for cold or archive storage?

### For Cluster Admin,

- How to install/upgrade each storage type (ceph/portworx) easily?
- How to migrate data from one storage to another type?
- Can storage layer provide better data governance or lineage information?



## Challenges in Storage Side: Multiple Cloud Storage Access

• •/•

Few open use cases are,

- Spark users often access data from s3, abfs, gfs for various I/O requirements.
- Could there be challenges in accessing data policies for each clouder vendors?

A uniform layer could be easy for Spark or other big-data deployment to access data from various cloud storages.

Why Spark on K8s

Compute: Gaps of K8s Scheduling for Spark

Spark Scheduling with YuniKorn & Features Overview

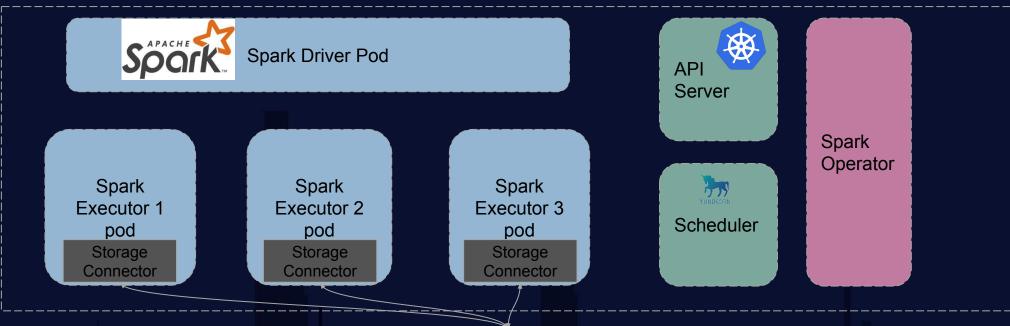
Storage: Challenges in K8s

**Common Storage Layer - Introduction** 



## SPARK ON K8S for Multi Cloud I/O Access





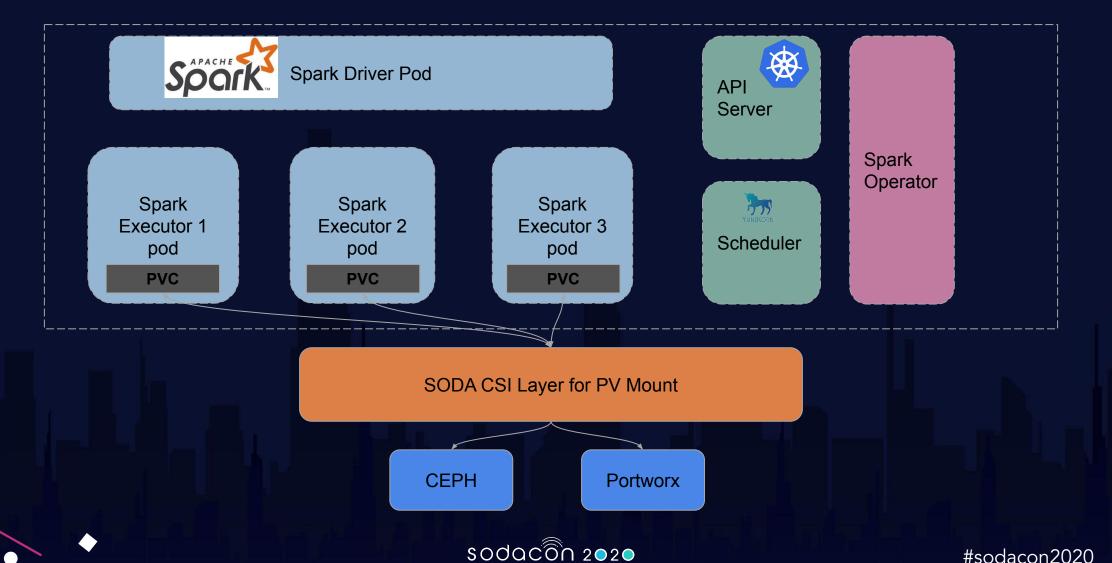
**SODA Layer** 

S3

**ABFS** 



## SPARK ON K8S for Persistent Volumes - Proposal





## **VISION - COMPUTE SIDE**



### Computes

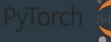
Data Engineering, Realtime Streaming, Machine Learning













### **Types**

Micro services, batch jobs, long running workloads, interactive sessions, model serving







### **Targets**

Multi-tenancy, SLA, Resource Utilization, Cost Mgmt, Budget









Unified Compute Platform for Big Data & ML



# SOCION 2020 DATA CONNECTED

#sodacon2020

THANK YOU!

