sodacon
— Global 2021 — July 13-14

Multi-Cloud data management in real-world usage

Paul Speciale

Chief Product Officer, Scality

soda foundation

Scality Multi-Cloud Data Management Solution

- Scality introduced Zenko as an Apache 2.0 open-source project in 2017 to help solve emerging problems in multi-cloud data management - on www.zenko.io and on https://github.com/scality/Zenko
- Now ~4 years of practical multi-cloud deployment experience with customers

Cloud data management in practical real-world use

- The 3 key use-cases
  - Cloud archiving
  - Cloud data bursting
  - Cloud Disaster Recovery (D/R)

- Commercial usage
  - DataOps / Analytics cloud service
  - Data storage solutions

Customer focus: 3 key solutions for cloud data

On-premises to cloud archive

- Offload cold/dormant data to cloud long-term archives (Azure Cold Storage, AWS Glacier or other)
- Solution: Lifecycle tiering policies - move data on data age/date rules

On-premises to cloud for compute bursting

- Move data to cloud compute services: AI/ML, analytics, transcoding, CDN, Video Indexing & others
- Solution: replication or tiering policies - based on tags/filters

On-premises to cloud for Disaster Recovery (D/R)

- Maintain a copy data in cloud for disaster recovery & business continuance in the event of site failure or outage
- Solution: continuous replication to cloud(s) to two copies with low RPO

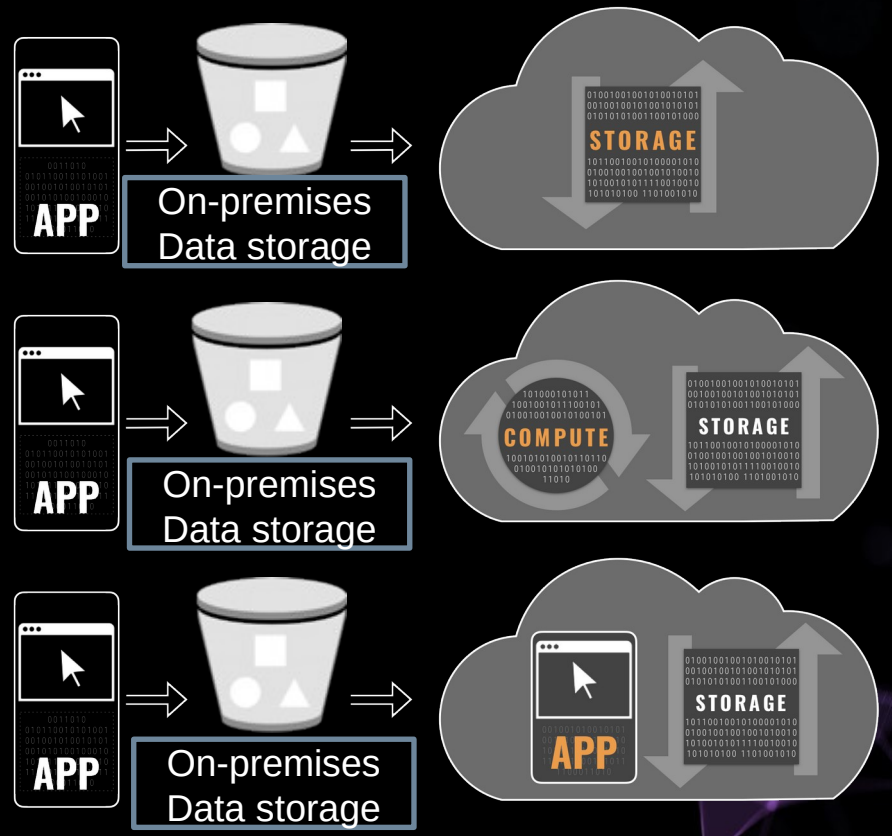

CLOUD D/R IN LIFE SCIENCES (BIOPHARMA RESEARCH)

Major US Pharmaceuticals Company

BIG DATA FOR BIOPHARMA GENOMICS DATA

- high-performance flash file system solution Weka.io - physical on-premises and virtual in AWS EC2
- 10PB object store (Scality) for on-premises persistent storage

BUSINESS CONTINUANCE- ZENKO FOR CLOUD D/R

- Replicates data from RING to AWS S3 via “out-of-band” updates
- Data failover to cloud after primary site outage & manage failback once primary site is restored

CUSTOMER VALUE FOR CLOUD D/R

- Time to market vs 2nd physical DC
- Reduced Capex vs. second full physical deployment

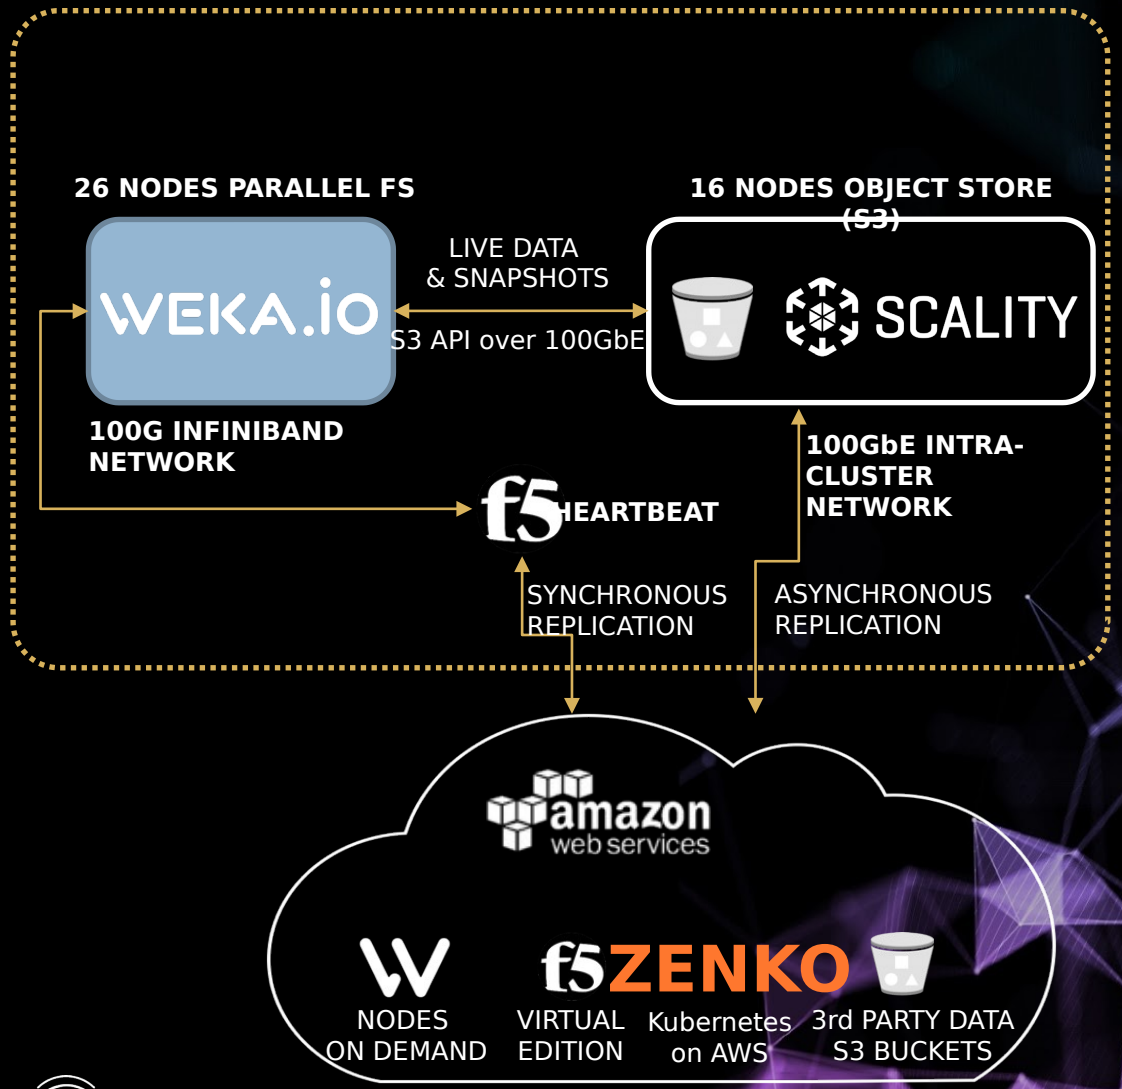

Bloomberg

The challenge: Multi-cloud data bursting in broadcast

A PREMIUM 24/7 MEDIA NETWORK

Bloomberg Media is the world's largest business news organization, reaching a premium audience of 91 million. The company delivers business and financial news across every platform and every time zone.

THE TECHNICAL CHALLENGE

To build a global system for file, object and cloud storage that would serve all of Bloomberg's storage and playback needs worldwide was a challenge:

- Cloud egress is costly and must be managed
- Storing content on multiple cloud providers has many benefits if managed properly.
- Public cloud storage doesn't guarantee 100% uptime.
- Global content lifecycle policies enable flexible, standardized and scalable data management.

Bloomberg

“We have made very good progress in standing up not only a product, but really a whole infrastructure solution that can add the scale we need for all the changes coming in the future.”
– Bloomberg Cloud CTO

The multi-cloud CDN solution

OBJECT STORAGE FOR MEDIA ASSETS
- Data centers in London & NYC

ZENKO: MEDIA WORKFLOWS TO MULTIPLE PUBLIC CLOUDS
- Replicates media files from on-premises to public clouds for transcoding & CDN services
- Simultaneous transfers “in” to 15+ regions in AWS, Azure & Google
- Avoids egress charges to move data between clouds

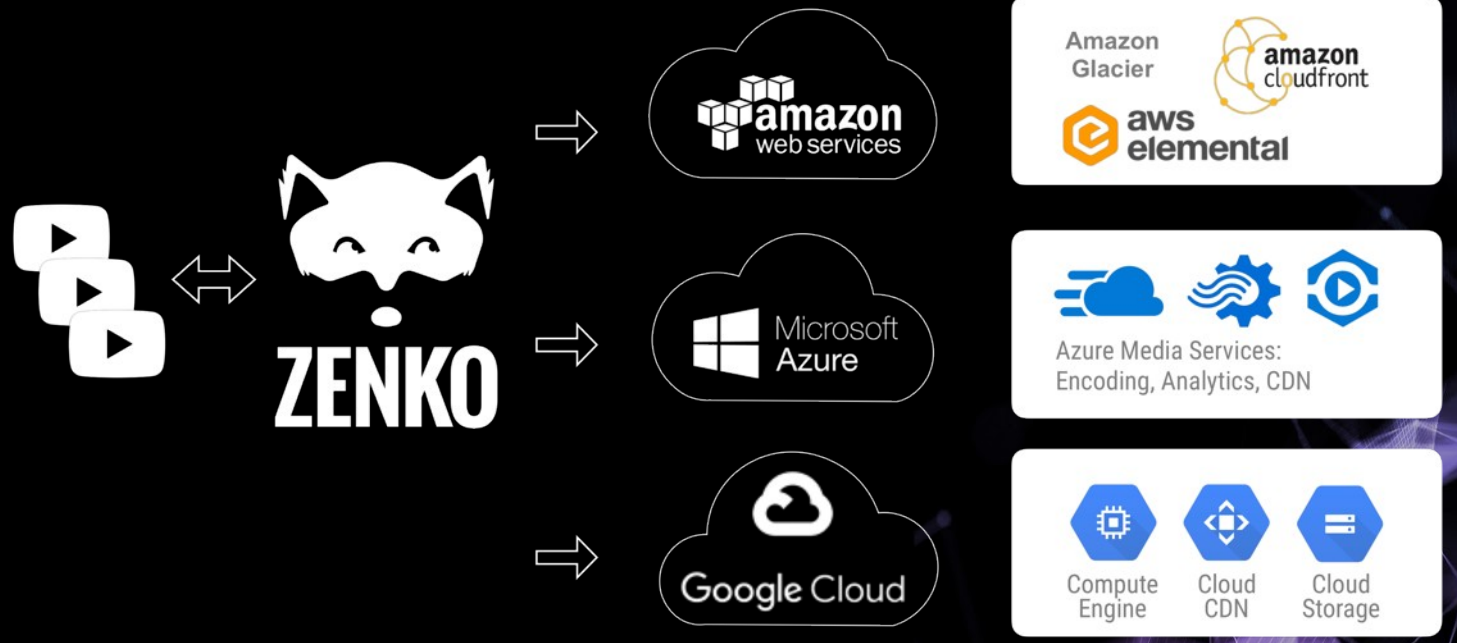

X Bloomberg Media Group

Commercial use of Zenko based multi-cloud solutions

- Saagie Data Orchestrator
  - Cloud-based DataOps & Analytics Platform
  - ETL & visualize data workflows across clouds
- Fujifilm Object Archive
  - S3 object to tape solution
  - Extended Zenko with AWS Glacier API
- Object Storage
  - Scality's own ARTESCA lightweight, cloud-native object storage solution

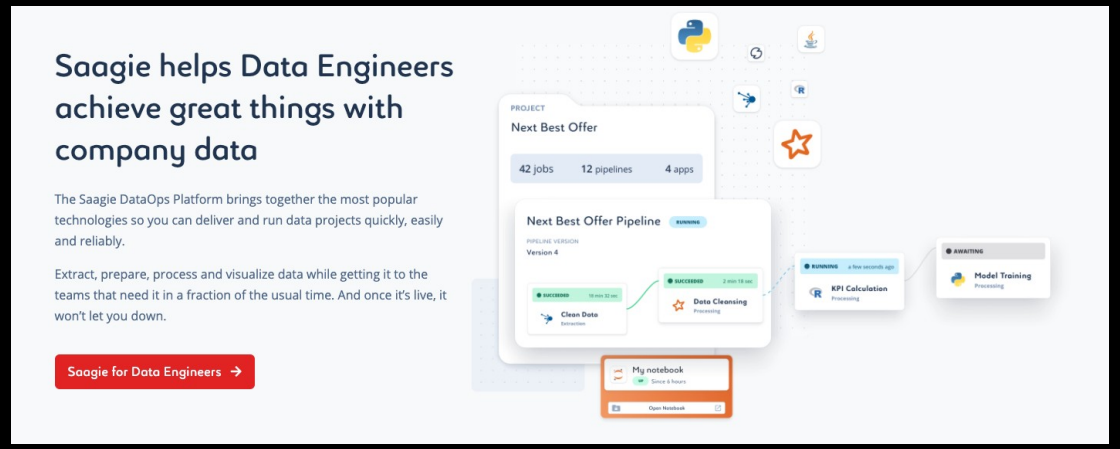

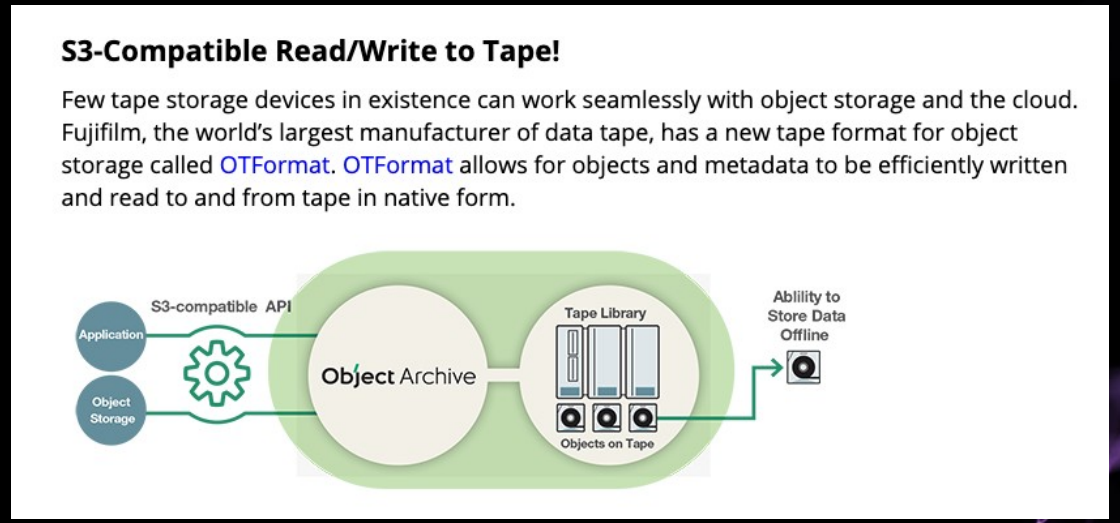

Summary

- Customers are solving real-world problems using multi-cloud data management
- 3 core use-cases have been the focus for the last 2-3 years
- New requirements are emerging for advanced/user-defined data management policies
  - Please join our 10-minute lightning talk session this afternoon for an introduction
  - **User-Defined Cloud Workflow Policies: today 4:05 pm - 4:15 pm PDT.**

sodacon
— Global 2021 — July 13-14

Thank You

Paul Speciale

Chief Product Officer, Scality
paul.speciale@scality.com

soda foundation